



(a) T1 · Communication / compute overlap**(b) T2 · Execution organization**

8 × NVIDIA H800 · 128K total tokens · rank-max p50

T1 Comm + compute = sum of isolated p50s. T2 uses preloaded K/V (compute-only).