


T3 · Transformer layer time by placement

All-CP BR-PBS hybrid Megatron-adapted Zeppelin-adapted

OtherBwd OtherFwd CoreFwd CoreBwd



8 × NVIDIA H800 · equal-case means (n=20 per placement) · all placements use the same MegaRing executor
Each iteration uses the rank with maximum full-layer CUDA FWD+BWD time; Others = full phase – core attention.